

Yiqing Xie

Carnegie Mellon University – 5000 Forbes Ave, Pittsburgh, PA 15213, USA
+1 (518) 763-5018 · yiqingxi@andrew.cmu.edu · yiqingxyq.github.io

Education

- Carnegie Mellon University** 2022 – 2027 (Expected)
- Ph.D. in Language and Information Technology
 - Advisors: [Carolyn Rosé](#) and [Daniel Fried](#)
 - Thesis: Synthetic Training Data Construction for Coding Agents
Committee: [Carolyn Rosé](#), [Daniel Fried](#), [Graham Neubig](#), and [Dan Roth](#) (UPenn)
- University of Illinois Urbana-Champaign** 2020 – 2022
- M.S. in Computer Science (GPA: 4.0/4.0)
 - Advisor: [Jiawei Han](#)
- Hong Kong University of Science and Technology** 2016 – 2020
- B.Sc. in Computer Science; double major in Mathematics (GPA: 3.9/4.3)

Research Highlights

- Agentic Post-Training.** I construct scalable and effective coding agent training environments [[COLM 2025](#)] and design lightweight training tasks that generalize to multiple coding agent benchmarks [[ICML 2026](#)].
- Coding Agents.** I develop multi-agent collaboration settings that are effective, flexible, and cost-efficient [[EMNLP 2026](#)] and develop automatic analysis methods that efficiently discover the influential factors of coding agent training demonstrations [[preprint](#)].
- Factuality and Trustworthiness.** I train critique-based evaluators that substantially improve factuality training [[ACL 2025](#)] and design evaluation frameworks for reliable and trustworthy clinical text generation [[ACL 2024](#)].

Industrial Experiences

- Meta AI** May 2024 – Oct. 2024
- Research Intern, GenAI
 - Manager: [Hejia Zhang](#); Peers: [Di Jin](#), [Pradyot Prakash](#), and [Sinong Wang](#).
 - Trained a fine-grained critique-based evaluator and used it to improve generators' factuality ([ACL 2025](#)).
- Microsoft Research, Redmond** June 2023 – Aug. 2023
- Research Intern, Health Futures
 - Managers: [Sheng Zhang](#), [Hao Cheng](#), and [Hoifung Poon](#).
 - Developed a multi-aspect, fine-grained evaluation framework for medical text generation ([ACL 2024](#)).
- Microsoft Research, Redmond** May 2022 – Aug. 2022
- Research Intern, Productivity and Intelligence Group
 - Manager: [Chenyan Xiong](#).
 - Worked on continuously pretrained models for zero-shot dense retrieval ([SIGIR 2023](#)).
- Alibaba DAMO Academy** July 2020 – Feb. 2021
- Research Intern, Data Analytics and Intelligence Lab
 - Managers: [Yaliang Li](#) and [Bolin Ding](#).
 - Worked on few-shot interaction recommendation across multiple scenarios ([WWW 2022](#)).

Selected Publications

For the full list, please refer to Google Scholar: scholar.google.com/citations?user=ov9EcNEAAAAAJ.

* denotes equal contribution.

Steer, Don't Solve: Training Small Critic Models for Large Code Agents EMNLP, 2026

Shubham Gandhi*, **Yiqing Xie***, Atharva Naik, Ruichen Zhu, and Carolyn Rose.

Mind the Sim2Real Gap in User Simulation for Agentic Tasks COLM, 2026

Xuhui Zhou, Weiwei Sun, Qianou Ma, **Yiqing Xie**, Jiarui Liu, Weihua Du, Sean Welleck, Yiming Yang, Graham Neubig, Sherry Tongshuang Wu, and Maarten Sap.

Hybrid-Gym: Training Coding Agents to Generalize Across Tasks ICML, 2026

Yiqing Xie, Emmy Liu, Gaokai Zhang, Nachiket Kotalwar, Shubham Gandhi, Sathwik Acharya, Xingyao Wang, Carolyn Rose, Graham Neubig, and Daniel Fried.

RepoST: Scalable Repository-Level Coding Environment Construction with Sandbox Testing COLM, 2025

Yiqing Xie, Alex Xie, Divyanshu Sheth, Pengfei Liu, Daniel Fried, and Carolyn Rose.

TheAgentCompany: Benchmarking LLM Agents on Consequential Real World Tasks NeurIPS, 2025

Frank F. Xu, Yufan Song, Boxuan Li, Yuxuan Tang, Kritanjali Jain, Mengxue Bao, Zora Z. Wang, Xuhui Zhou, Zhitong Guo, Murong Cao, Mingyang Yang, Hao Yang Lu, Amaad Martin, Zhe Su, Leander Maben, Raj Mehta, Wayne Chi, Lawrence Jang, **Yiqing Xie**, Shuyan Zhou, and Graham Neubig.

Improving Model Factuality with Fine-grained Critique-based Evaluator ACL, 2025

Yiqing Xie, Wenxuan Zhou, Pradyot Prakash, Di Jin, Yuning Mao, Quintin Fettes, Arya Talebzadeh, Sinong Wang, Han Fang, Carolyn Rose, Daniel Fried, and Hejia Zhang.

CodeRAG-Bench: Can Retrieval Augment Code Generation? Findings of NAACL, 2025

Zora Zhiruo Wang, Akari Asai, Xinyan Velocity Yu, Frank F. Xu, **Yiqing Xie**, Graham Neubig, and Daniel Fried.

DocLens: Multi-aspect Fine-grained Evaluation for Medical Text Generation ACL, 2024

Yiqing Xie, Sheng Zhang, Hao Cheng, Pengfei Liu, Zelalem Gero, Cliff Wong, Tristan Naumann, Hoifung Poon, and Carolyn Rose.

Data Augmentation for Code Translation with Comparable Corpora and Multiple Findings of EMNLP, 2023

References

Yiqing Xie, Atharva Naik, Daniel Fried, and Carolyn Rose.

Unsupervised Dense Retrieval Training with Web Anchors SIGIR, 2023

Yiqing Xie, Xiao Liu, and Chenyan Xiong.

KoMen: Domain Knowledge-Guided Few-Shot Interaction Recommendation on Multiplex Networks WWW, 2022

Yiqing Xie, Zhen Wang, Carl Yang, Yaliang Li, Hongbo Deng, Bolin Ding, and Jiawei Han.

Invited Talks

Designing Scalable and Generalizable Synthetic Training Tasks for Coding Agents

◦ Nvidia Reading Club, Nvidia June 2026

Synthetic Data Approaches for Coding Agent Training

◦ Guest Lecture for 11-891, Neural code generation, Carnegie Mellon University Nov 2025

Constructing Scalable Repository-Level Environments for Coding Agent Training

◦ EFML Reading Group, Stanford / University of Washington Apr 2025

Multi-aspect Fine-grained Evaluation for Medical Text Generation

◦ Health and Life Science Seminar, Microsoft Research Apr 2024

Honors and Awards

◦ CMU Presidential Fellowship in LTI 2024 – 2025

◦ Siebel Scholar, Class of 2022 2021 – 2022

◦ HKUST Academic Achievement Medal (top 1%) 2020

◦ HKSAR Government Scholarship Fund – Reaching Out Award 2018

◦ HKUST Scholarship for Continuing Undergraduate Students 2017 – 2019

◦ Dean's List, HKUST (three times) 2017 – 2019